

Dhrumil Ankola

(951) 213-8774 | ankoladhrumil@gmail.com | [linkedin.com/in/dhrumil-ankola](https://www.linkedin.com/in/dhrumil-ankola) | github.com/dhrumilankola | dhrumilankola.netlify.app

Summary

M.S. in Computer Science with deep expertise in designing and deploying production-grade AI systems and scalable data platforms. Proven track record of architecting solutions that deliver significant business impact—from boosting user retention by 40% to increasing industrial plant throughput by 15%. Expert in the full development lifecycle, encompassing LLM/RAG pipelines, microservices on Kubernetes, and large-scale data processing with Spark and Airflow

Experience

Shifa Precision Inc

06/2025 - Present

Software Engineer - AI/ML

Boston, MA

- Owning the backend development for Project Oasis, a proactive wellness engine that uses a Neo4j knowledge graph on AWS to generate patient "digital twins"
- Engineering a scalable data pipeline to unify biomedical data (ClinVar, PubMed, openFDA) into a knowledge graph with over 1M+ nodes and 10M+ relationships
- Building an AI-driven backend featuring LLMs for automated data extraction and a novel Graph Neural Network (GNN) framework for predictive health analysis

Center for Environmental Research and Technology

07/2024 - 04/2025

Assistant Software Developer

Riverside, CA

- Engineered a full-stack inspection platform (React, Laravel) for the California Air Resources Board (CARB), improving page-load times by 30% while serving 5,000+ monthly users
- Automated the regulatory data pipeline using AWS S3 and Glue, which enhanced data accuracy by 25% and enabled near-real-time reporting capabilities
- Implemented a Jenkins-driven CI/CD pipeline and authored comprehensive PHPUnit test suites, reducing deployment rollbacks by 30% for a system processing over 1 TB of data monthly

Edicon Paper

01/2022 - 07/2023

Product & Systems Engineer

Gujarat, India

- Architected a full-stack B2B client portal (React, Node.js, PostgreSQL) to manage high-volume orders for a 150,000-ton/year facility, cutting client service inquiries by 30%
- Engineered a data pipeline (Python, Airflow, Spark) analyzing real-time sensor data to optimize machine scheduling, which increased overall plant throughput by 15%
- Deployed a quality assurance monitoring system (Python, InfluxDB) on a 400+ ton/day production line, identifying key defects to cut material waste by 5% and save over \$50,000 annually

Projects

LLM-Powered Anomaly Detection for Microservices Reliability

GKE, Kubernetes, Docker, OpenTelemetry, Vertex AI (RAG), Go, Python

- Engineered a scalable polyglot microservices platform on GKE with a zero-downtime CI/CD pipeline
- Built an LLM-powered anomaly detection engine using PEFT (F1 \approx 0.90) and integrated Vertex AI RAG for contextual alerting; boosted detection precision and automated root cause insights, cutting resolution time by ~40%

Study Buddy - Berkeley CalHacks

React, Django, LangChain, RAG, Gemma 2, ChromaDB, ElevenLabs AI

- Built AI tutoring platform using Elevenlabs conversation AI and fine-tuned Gemma 2 (4-bit LoRA) for 90% response accuracy
- Integrated RAG with ChromaDB for context-aware retrieval; enabled interactive mind maps and voice-driven aids

PyAsterix - Database Driver for AsterixDB

Python, Pandas, AsterixDB, NoSQL, SQL++

- Developed PyAsterix Python driver for AsterixDB with Pandas integration and SQL-style queries
- Tuned connection pooling and serialization to cut query latency by 15% on 10 M+ record workloads
- Implemented AsterixDataFrame API mirroring Pandas, simplifying distributed filtering and aggregation

Education

University of California, Riverside

03/2025

Master of Science in Computer Science (GPA: 3.7)

BVM Engineering College, Anand, India

07/2023

B.Tech in Information Technology (GPA: 3.8)

Skills

Languages: Python, Go, JavaScript, C++, SQL, NoSQL

Cloud & DevOps: AWS (S3, Glue, Lambda), GCP (GKE, Vertex AI), Kubernetes, Docker, Jenkins, OpenTelemetry, Git

AI/ML: PyTorch, TensorFlow, LangChain, RAG, Fine-Tuning (PEFT, LoRA), GNNs, LLMs (Gemma, Llama3, BERT)

Data Engineering & Databases: Apache Spark, Apache Airflow, PostgreSQL, MySQL, MongoDB, Neo4j (Graph), ChromaDB (Vector), InfluxDB (Time-Series), AsterixDB

Frameworks & Libraries: React.js, React Native, Node.js, Django, Laravel, Express.js

Architecture & APIs: Microservices, Distributed Systems, GraphQL, REST

Testing & Methodologies: PyTest, PHPUnit, Selenium, Agile/Scrum, Jira